

Be Named First

Operations Handbook

AI Search Visibility Consultancy

benamedfirst.com

Version 1.0 · August 2026

Internal reference document. Also intended as context for AI assistants.

Review quarterly — timelines, engine behaviour and pricing benchmarks shift.

Contents

Click any entry to jump to that section.

0	How to use this document
1	What the business is
2	Positioning and language
3	Why this market exists right now
4	The ideal client
5	Services and pricing
6	Expectations versus reality
7	How AI search actually works — technical foundation
8	The prompt set — the core instrument
9	Delivery — phase by phase, step by step
10	Reporting
11	Client acquisition
12	Sales and closing
13	Targets and unit economics
14	Risk register
15	Team, tools and delegation
16	First 30 days
17	Integrity rules — binding
18	Sources
19	Glossary
20	Quick reference card

0. How to use this document

This is the single source of truth for the business. It exists for two readers.

For you (human). Read Sections 1–7 once, end to end, before your first outreach. After that, use it as a reference: Section 9 while delivering, Section 11 while prospecting, Section 12 on sales calls.

For an AI assistant. Paste the relevant section as context before asking for work. Every section is written to be self-contained, so you can hand over one section without the rest and it will still make sense.

Useful prompts to pair with this document:

Task	Give the assistant
Write outreach emails	Sections 4, 11
Build an audit report	Sections 9, 10
Handle a sales objection	Sections 6, 7, 12
Qualify a prospect list	Section 4
Plan a client's fix roadmap	Sections 7, 9
Sanity-check a claim before saying it	Sections 6, 17

Rules for anything produced from this document

1. **Never promise a guarantee.** Not rankings, not positions, not citations, not timelines for specific outcomes. Section 6 explains why.
 2. **Never fabricate.** No invented client names, results, testimonials, logos, or statistics. Section 17 is binding.
 3. **Cite the source for every statistic.** If it is not in Section 18, do not state it as fact.
 4. **Distinguish what is verified from what is estimated.** Estimates in this document are labelled as such.
-

1. What the business is

One sentence

We get businesses named and cited inside AI-generated answers, and we measure it.

The longer version

When someone asks ChatGPT, Perplexity, Gemini, Claude, or Google's AI Overviews for a recommendation, they do not get a list of ten blue links. They get a written answer that names three or four businesses. Everyone else is invisible — not ranked low, **absent**.

Most businesses do not know this is happening to them. There is no dashboard that tells them. There is no Search Console for AI answers.

We do three things:

1. **Measure** where a business stands across the major AI engines, against named competitors, using a fixed and repeatable prompt set.
2. **Fix** what we find — on-site work so answers can be extracted from their content, off-site work so the engines trust and cite them.
3. **Report** movement monthly, per engine, with the same prompt set, so progress is provable rather than asserted.

Why we win business

Not because we are cheaper. Because of two things:

Proof at first contact. We do not pitch. We show the prospect a screenshot of an AI engine naming their competitors and not them. That is evidence, not marketing. It converts because it is specific and it stings.

Measurement as the product. Most vendors in this space sell vibes — "AI-optimised content" with nothing behind it. We run a fixed prompt set, repeatedly, under controlled conditions, and report numbers. Clients who care about measurement are the clients who stay.

What we are not

- Not an SEO agency. Related discipline, different mechanics, different metrics.
 - Not a content mill.
 - Not a tool vendor. We are a service. Tools may support us; they are not the product.
 - Not a guarantee business. See Section 6.
-

2. Positioning and language

The two disciplines we deliver

Clients do not need to know these terms. You do, because they map to different work, different timelines, and different reporting.

AEO — Answer Engine Optimization. Making the business's content the thing an engine **extracts** as a direct answer. Surfaces: featured snippets, People Also Ask, voice assistants, Google AI Overviews. Mostly **on-site** work. Faster to move.

GEO — Generative Engine Optimization. Making the business a **cited source** when an LLM synthesises an answer. Surfaces: ChatGPT, Claude, Perplexity, Gemini. Heavily **off-site** work — third-party mentions, reviews, original data, entity consistency. Slower, more defensible.

	AEO	GEO
Mechanism	Extraction — engine lifts your text	Synthesis — engine blends sources, you are one
Where the work is	Mostly on-site	Mostly off-site
Winner	Usually one page wins	Multiple sources blend
Metric	Did we get extracted — binary, per query	Citation share across many prompts — probabilistic
Speed	Faster	Slower

Simplest framing: AEO wants you to *be* the answer. GEO wants you to be in the set of sources the answer was *built from*.

What to say to clients

Sell one thing: **AI Search Visibility**. Do not lead with acronyms.

The client's actual question is always the same: "*Why doesn't my business come up when someone asks ChatGPT?*" Answer that.

Explain the AEO/GEO distinction only when the buyer is technical. Then it becomes a credibility asset, because most vendors cannot articulate the difference.

Language rules

Use: visible, cited, named, measured, baseline, prompt set, citation share, mention rate.

Avoid: guaranteed, ranked #1, dominate, unlock, revolutionary, game-changing, 10x, hack, secret. Also avoid "AI-powered" as a description of our service — we are a service applied to AI systems, not an AI product.

3. Why this market exists right now

Facts, with sources. Use these in outreach and on sales calls. Do not embellish them.

The shift is measurable

- AI search visits grew **42.8% year over year** while Google grew **2.4%**, and AI-referred visitors are estimated at roughly **4.4x** the value of a regular organic visitor. [3way]
- Cited links in AI overviews earn click-through rates of roughly **8–12%**, while uncited competitors are absent from the conversation entirely, per research from Princeton and Georgia Tech. [SomethingInc]
- **97%** of surveyed practitioners reported a positive impact from AEO in 2025, and **94%** planned to increase AEO investment in 2026. [Conductor]
- Searches for "geo agency" are up **2,300% year over year**. [LLMPulse]
- The Generative Engine Optimization market is forecast at **USD 1.09B in 2026**, growing to **USD 17.1B by 2034** at a 40.6% CAGR. [Algomizer]

The money is already allocated

- Credible AEO/GEO retainers for mid-market businesses land at **\$2,000–10,000 per month**. One-off audits run **\$1,000–4,000**. Enterprise programmes reach **\$25,000+**. [OutreachBloom, Digital Elevator]
- Full AEO/GEO programmes run by specialist agencies reach **\$8,000–15,000 per month**. [Fuel Online]

This matters for sales: **we are not asking for a new budget**. We are asking to redirect an existing SEO or marketing budget toward a channel the client is already losing.

Building websites is not the business

Relevant because it explains why we sell this and not web development:

- In a 2026 survey, **75–76%** of web design professionals said AI competition had already affected their business, and ranked it a bigger threat than shrinking budgets. [20i, TechRadar]
- AI builders have collapsed the price floor for a professional business site from thousands to under a few hundred dollars. [pikapod analysis]

Where a client can compare you to a chatbot, you have no pricing power. AI search visibility is not yet commoditised. That window is why this business exists now.

All sources listed in Section 18. Content rephrased for compliance with licensing restrictions.

4. The ideal client

This is the most important operational section. A wrong client costs three months and ends in a refund request.

The three-part test

A prospect must pass all three.

- 1. A new customer is worth \$2,000 or more to them.** If one new customer is worth \$5,000, a \$1,500 audit pays for itself with a single conversion. If a customer is worth \$50, you will never justify the fee. This single filter matters more than every other consideration combined.
- 2. They already spend money on marketing.** You are redirecting a budget, not creating one. Businesses with no marketing spend have no mental category for this purchase.
- 3. Their competitor is cited and they are not.** This is the emotional trigger. No gap, no urgency, no sale. If a prospect is already well cited, they are not a prospect — they are a reference case.

Target verticals

Ranked by fit. Pick **one** and stay in it until you have three case studies.

Vertical	Why it fits
Cosmetic and implant dentistry	Very high per-patient value, competitive local markets, owner-operated
Med spas / aesthetics	High ticket, high growth, marketing-literate buyers
Elective and private healthcare	High value, trust-driven purchase, AI-researched by patients
Personal injury / immigration law	Extremely high case value, already heavy marketing spenders
B2B SaaS	Understands the problem instantly, sophisticated on measurement
Wealth / insurance advisory	High lifetime value, credibility-led purchase
High-value home services (solar, roofing)	Large ticket sizes, local competition

Avoid: restaurants, retail, low-ticket local services, anything with a sub-\$500 customer value, and anything with no digital dependency.

One vertical or several — the strategy

Two decisions get conflated here. Separate them.

Decision	Answer
Positioning — what the site, LinkedIn and pitch say	One vertical. Always.
Prospecting — who you actually email	One primary + one adjacent test, first 60 days only

Showing multiple verticals in your positioning makes you a generalist. Generalists discount; specialists command a premium.

Why one vertical compounds:

- **Prompt set reuse.** A new vertical costs 2–3 hours of prompt-set work. A new client in an existing vertical costs 30–45 minutes. Build once, reuse forever.
- **Off-site knowledge does not transfer.** The directories, review platforms and industry publications that engines pull from are completely different for dental versus legal. After five clients in one vertical you know exactly which six sources matter. That knowledge is your margin.
- **Case studies only convert inside a vertical.** "Three dental clinics" converts a dentist. "One dentist, one lawyer, one SaaS" converts nobody.
- **Referrals only travel inside a vertical.** Dentists know dentists. This is how outbound eventually stops.
- **Pricing power comes from specificity.** "AI search visibility for cosmetic dentistry" prices better than "AI visibility for businesses".

Why not to bet 100% on one either: you have no data yet. Committing entirely to one vertical and discovering its owners do not reply costs four weeks.

The hedge: 40 touches in the primary vertical, 15–20 in an adjacent backup. Compare reply rates. At 60 days, kill the loser and go all in on the winner. This is cheap because you adapt one prompt set rather than building two.

The backup must be adjacent. Adjacency is what keeps the transfer cost low.

- **Good pair:** cosmetic/implant dentistry ↔ med spa / aesthetics. Same buyer psychology (elective, high ticket, local, owner-operated), overlapping directories, similar prompt shapes. Roughly 60–70% of the prompt set carries over.
- **Bad pair:** cosmetic dentistry ↔ B2B SaaS. Different buyer, sales cycle, prompts, off-site levers and pricing. That is two businesses, not one.

Expand geography, not vertical

This is the rule most people get backwards.

When the dental clinics in Dubai are exhausted, move to Abu Dhabi, then Riyadh, then London, then Toronto. Same prompt set, same off-site playbook, same case studies, same accumulated pattern recognition. Only the city changes.

Changing vertical resets you to zero. Changing geography keeps everything.

Add a second vertical only when all three are true:

1. Three case studies exist in the current vertical
2. Referrals are arriving unprompted
3. Three or more cities have been worked through

Realistically that is 8–12 months out.

Current selection

Primary: cosmetic and implant dentistry, Dubai. Highest per-patient value, ~1.5 hour timezone gap, owner-operated single decision makers, English-language business environment, and a purchase buyers actively research through AI.

Backup test: med spas / aesthetics, Dubai. Adjacent, so the prompt set and off-site knowledge transfer.

Review date: 60 days from first outreach. Kill the weaker performer, then scale by city.

Geography

Pricing forces Western and Gulf markets. An Indian SMB will not pay ₹1.7L/month for a service it does not yet understand.

Market	Assessment
UAE / GCC	Best first target. High-ticket clinics and firms, Western-level rates, English-language business environment, and a timezone gap of only ~1.5 hours from India — meaning same-day conversation instead of next-day.
US	Largest market, most budget. Timezone gap of 9–12 hours makes iteration slow.
UK / EU	Good rates, manageable timezone, EEA has raised general digital-compliance awareness.
Canada / Australia	Viable, smaller volume.
India (B2B SaaS selling globally)	Viable secondary. Their buyers are Western, so the problem is real. Lower price point, but far easier trust and same timezone. Realistic source of a first client.

Stage 1 qualification — desk research, 5 minutes per prospect

Run this **before** any contact. It saves more time than anything else in this handbook.

#	Check	How to verify	Fail means
1	Customer value \$2,000+	Services and prices on their site	Skip — fee never justifiable
2	Already spending	Google Ads running, agency credit in footer, recent blog activity, marketing roles posted on LinkedIn	Skip — no budget category
3	Visibility gap exists	Run 3–5 prompts from the vertical prompt set	Skip — no pain, no sale
4	Site is modifiable	Identifiable CMS, or a developer they can reach	Skip — cannot deliver AEO
5	Decision maker reachable	Owner-operated, or small enough that the owner decides	Skip — procurement will kill a \$1,500 deal
6	Not already served	No in-house SEO team, no large incumbent agency	Redirect — approach the agency for white-label work instead

Scoring: 4 or more passes → add to outreach list. 3 or fewer → skip.

Check 6 is frequently misplayed. Where an agency already holds the account, the client is not your prospect — the agency is. They have the client relationship and no AI-visibility capability. That is a partnership, not a competition.

Stage 2 qualification — on the call

Six questions. Ask them early; disqualify fast.

1. **"What is a new client worth to you over their lifetime?"** Does not know, or the number is small → out.
2. **"What are you spending on marketing now, and with whom?"** Confirms budget exists and reveals the incumbent.
3. **"Who can make changes to your website, and how quickly?"** Delivery feasibility. A locked platform is a hard blocker.
4. **"If this works, who signs off on continuing?"** Decision authority.
5. **"What would you need to see in 90 days to keep going?"** Sets expectations and exposes impossible demands early.
6. **"Have you tried anything like this before?"** Previously burned is harder to sell but often more sophisticated.

Red flags — walk away

- Asks for a **guarantee**, or "pay on results". The single biggest flag. LLM outputs are non-deterministic; a guarantee would be a lie.
- Wants a promised ranking or position.
- Expects complete results within weeks.
- Negotiates the audit price down hard on the first call.
- Needs three people to approve a \$1,500 decision.
- Site is on a locked platform they cannot modify.
- Business is entirely offline walk-in with no digital dependency.

Green flags — prioritise

- Reacts **emotionally** to the competitor screenshot. Strongest buying signal there is.
- Already aware AI search matters, and worried about it.
- Asks **"how will you measure this?"** — sophisticated buyer, will value the differentiator, and will retain.
- Owner-operated, single decision maker.
- High-ticket service.

The discipline

You are one person. Your scarcest resource is time, not leads. Disqualifying is not lost revenue — it is protected capacity.

5. Services and pricing

Three tiers, designed as a ladder. Most clients enter at the audit.

Tier 1 — AI Visibility Audit

Price: \$1,500 fixed. (First three clients: \$750 as a founding-client rate, in exchange for a written case study. See "Founding client rate" below.)

Payment: 100% upfront. Normal for fixed-scope diagnostic work.

Duration: 2 weeks from prompt-set sign-off.

Deliverables:

- Baseline mention rate and citation share across the engines in scope, per engine, against up to three named competitors
- Diagnosis across four areas: crawler access, extractability, structured data, off-site footprint
- A prioritised fix roadmap, ordered by expected movement
- The prompt set itself, documented and versioned

Explicitly in scope: measurement, diagnosis, recommendations. **Explicitly out of scope:** implementing any of the fixes.

The roadmap is theirs to keep. They can execute it themselves or hand it to another vendor. Say this out loud — it lowers perceived risk and signals confidence.

Tier 2 — Implementation

Price: from \$5,000 per project. Scoped from the audit roadmap.

Payment: 50% on start, 50% on delivery.

Duration: 4–6 weeks for the sprint. Off-site work continues beyond it.

Deliverables:

- Key pages restructured into extractable, answer-first content
- Structured data (JSON-LD) implemented and validated
- Crawler access configured and verified in server logs
- Entity consistency corrected across the web
- Off-site work initiated: directories, review platforms, third-party placements, original data asset

Tier 3 — Monthly Visibility Monitoring

Split into two, deliberately. This is a change from the original single tier and the reasoning matters.

3a — Measurement retainer: from \$1,500/month.

- Fixed prompt set re-run monthly under controlled conditions

- Mention rate and citation share, per engine, versus named competitors
- Trend reporting against baseline
- A prioritised fix list — delivered, not executed

Nearly pure margin. Low delivery load. **This is the churn defence** — the client sees numbers every month even before citations move.

3b — Measurement + execution: from \$3,000/month.

Everything in 3a, plus the fix list actually executed each month.

Minimum term: 3 months. State the reason honestly: the GEO validation cycle is itself 3–6 months. [CNABKE]
Selling month-to-month guarantees churn at month 2, before results can land.

Why the original single monitoring tier was wrong

A published range of \$2,000–10,000/month is a 5x spread. That tells a buyer two things: the vendor does not know what this costs, and they should anchor at the bottom. Splitting measurement from execution removes the spread and creates a genuinely better entry product.

Public pricing policy

Show pricing on the site, with this structure:

Tier	Display
Audit	Fixed: \$1,500
Implementation	From \$5,000
Monitoring	From \$1,500/month

Fixed scope gets a fixed number. Variable scope gets "from". Never publish a wide range — it is worse than publishing nothing.

Why public pricing helps this specific business: the channel is outbound, so the sequence is cold email → prospect checks the site → call. The site's job is credibility and qualification, not conversion. A published price filters out prospects with no budget before they consume a call slot. For a solo operator that is the largest single benefit. It also differentiates: this niche is full of vendors who will not name a number.

Never publish: hourly rates, wide ranges, any guarantee, or client counts that are not true.

Founding client rate

First three clients: **\$750 audit**, in exchange for a written, publishable case study.

Position it as an exchange, not a discount: *"Founding client rate — I'm building public case studies in this vertical, and I'll trade half the fee for permission to publish the results."*

This only works because the public price is \$1,500. A visible anchor is what makes a concession valuable. Once three case studies exist, the rate goes to \$1,500 and stays there.

ROI framing — use this next to every price

Prices sell badly on their own. Attach the client's own economics:

"A new implant patient is typically worth \$3,000–5,000 to your practice. This audit costs less than half of one patient."

"Your average case value is \$15,000. If this produces one additional case in a year, the retainer has paid for itself four times over."

This does more for conversion than any change to the number itself.

6. Expectations versus reality

The most important section for avoiding refunds. Read before every sales call.

The core mechanic clients do not know

AI search runs on **two separate memory systems**, and only one of them is reachable.

1. The retrieval layer — reachable. Perplexity runs a live web search on essentially every query. ChatGPT Search and Google AI Overviews also pull live content. A new authoritative source can begin shaping answers within **hours to weeks**. [FiveBlocks, Forrester]

2. The training layer — not reachable. A model's baseline knowledge only changes when the model is retrained or fine-tuned, a cycle measured in **months**. Claude anchors primarily to its training corpus, using Brave Search only when retrieval is triggered. [Browworks]

Everything we sell targets the retrieval layer. We never promise anything about the training layer. It is outside our control and outside anyone's.

Engines are not interchangeable

When the same prompt is run across ChatGPT, Perplexity, Gemini, and Google AI Overviews, the overlap in cited domains is roughly **11%**. Four platforms, four almost entirely separate source lists. [Topify]

Consequences:

- "We'll make you visible in AI" is a meaningless promise. Work and reporting are **per engine**.
- Six engines means roughly six times the measurement labour. Price accordingly.
- Wins arrive on different engines at different times.

One more useful data point: businesses with a detailed Wikipedia entry get their first ChatGPT citation in an average of **28 days**, versus **52 days** without one. Wikipedia alone contributes about **7.8%** of ChatGPT's citations and makes up roughly **48%** of its top-10 cited sources. [Topify]

What the client assumes versus what is true

Client assumes	Reality
"You can get me to #1 in ChatGPT"	There is no #1. Answers are generated and non-deterministic. The same prompt can return different answers on consecutive runs.
"This is just SEO"	Related but distinct. Google rankings and AI citations do not map one-to-one.
"I'll see results in two weeks"	Crawler access changes in days. First citations typically 4–12 weeks. Consistent visibility 3–6 months.
"You control what the AI says"	We influence which sources it can find and trust. We do not control output.
"One engine's result applies everywhere"	~11% citation overlap between engines.
"If it doesn't work I get my money back"	The audit is a diagnostic deliverable. It is complete when delivered. No results-based refund.

The honest timeline

Present this table on sales calls. Do not soften it.

Period	What happens	Verified by
Days 1–7	Crawler access and indexable answers corrected. Fixable in days.	LeanLabs
Weeks 1–2	First AI mentions may appear. Volatile. Signal, not a win.	Higoodie
Weeks 4–8	Early signals, impressions rising. First mentions commonly 30–60 days.	WebFX, AirOps
Weeks 4–12	Citations begin on specific, low-competition prompts.	LeanLabs
Months 3–6	Repeat citations, consistent visibility across engines.	WebFX, Upgrowth

The sentence to say: *"Month one is plumbing. Months two and three are when citations start. Months three to six are when visibility becomes consistent."*

Promise anything faster and you will be issuing refunds.

The churn problem, and the answer

Results take 3–6 months. Retainers churn at month 2 or 3 if the client sees nothing. That is the central operational risk in this business.

The answer is not to overpromise. It is to **report leading indicators in weeks 1–4** — things that move before citations do and are independently verifiable:

1. **Crawler hits from server logs.** The strongest week-one proof. A 30-day site log study found GPTBot to be the most aggressive crawler at roughly **4,200 hits per site per day**, ClaudeBot at ~1,800, PerplexityBot at ~980, and all four major bots respecting robots.txt. [Digital Applied] So when you unblock crawler access,

the effect appears in logs **the next day**. "These bots weren't reaching your site at all; now they're here daily" is a real, verifiable win in week one.

2. **Schema validation** — before and after, via Google's Rich Results Test.
3. **Prompt-set coverage** — a mention without a citation is still movement. Report it.
4. **Featured snippets and PAA** — these move faster than LLM citations. AEO's speed advantage.
5. **Third-party placements secured** — reportable the moment they go live, before any engine cites them.

This discipline is what keeps a client to month three.

What we never promise

- Any change to a model's training layer
- A specific ranking or position
- Identical results across engines
- A specific number of citations
- Any guarantee whatsoever

What we do promise

- A documented baseline before any work starts
- A fixed, written scope of work
- Transparent reporting on a fixed schedule, using the same prompt set and the same conditions every time
- That the audit roadmap is theirs regardless of whether they continue

We sell a **process**, executed rigorously, and honest measurement of what it produced. That is the only ethically defensible position in this niche, and it is also the one that retains clients.

7. How AI search actually works — technical foundation

Everything we do maps to something in this section. Read it before delivering.

The five surfaces

Engine	Retrieval behaviour	Implication for us
Perplexity	Live web search on essentially every query, continuously refreshed index, shows sources by design	Fastest to move. Win here first.
ChatGPT Search	Blends training knowledge with a web search layer. Browsing on → days. Browsing off → months.	Medium speed. Wikipedia presence matters here more than elsewhere.
Google AI Overviews	Tied to Google's index	Follows normal indexing cadence — weeks
Gemini	Google index-backed	Similar to AI Overviews
Claude	Anchors primarily to training corpus; Brave Search when retrieval triggers	Slowest. Set expectations accordingly.

Sources: Browworks, FiveBlocks, DataStudios, Forrester.

Engine sequencing strategy. Do not attack all engines at once. Win Perplexity first — it is the fastest and gives you a month-one reportable result. Then AI Overviews and ChatGPT. Broad consistency by month three.

What makes an engine cite a source

Three requirements, in order. Each is a failure point.

1. Access. The engine's crawler must be able to fetch the page. Blockers: `robots.txt` disallow, JavaScript-only content, cookie walls, interstitials, accidental `noindex`, login gates. If a crawler cannot read it, nothing else matters.

2. Extractability. The content must be structured so an answer can be lifted from it cleanly. Requirements: answer-first paragraphs, question-shaped headings, real semantic HTML, genuine `<table>` markup with `<th scope>`, self-contained statements with no "as mentioned above", valid structured data.

3. Trust. The engine must consider the source credible enough to name. Signals: independent third-party mentions, editorial coverage, reviews on trusted platforms, presence in knowledge bases, consistent entity information, original data worth citing, freshness.

Point 3 is where GEO lives and it is mostly **off-site**. This is why GEO is slower and why it is more defensible once achieved.

The AI crawlers

Configure access for all of these:

GPTBot	OpenAI, training
OAI-SearchBot	OpenAI, search index
ChatGPT-User	OpenAI, user-triggered fetch
ClaudeBot	Anthropic
anthropic-ai	Anthropic
PerplexityBot	Perplexity
Google-Extended	Google, Gemini/AI training control
Applebot-Extended	Apple
CCBot	Common Crawl (feeds many models)
Bingbot	Bing, feeds ChatGPT search

Observed crawl volumes from a 30-day site log study: GPTBot ~4,200 hits/site/day, ClaudeBot ~1,800, PerplexityBot ~980. All four major bots respected robots.txt. [Digital Applied]

Operational consequence: GPTBot crawls daily, so an access fix shows up in logs within 24–48 hours. That is your first reportable win on every engagement.

llms.txt

A plain-text file at the site root summarising who the business is, what it offers, and where the canonical pages are. Adoption is not universal across engines and its impact is not independently proven at the level schema is. Treat it as low-cost, low-risk, not as a primary lever. Do not oversell it to clients.

Structured data that matters

Schema type	Where	Why
Organization / LocalBusiness	Site-wide	Core entity definition
Industry-specific (Dentist , MedicalClinic , LegalService)	Site-wide	Category precision
Service	Each service page	What is offered
FAQPage	FAQ and any page with Q&A	Direct extraction target
Article + author + dates	Content pages	Authority and freshness
DefinedTerm / DefinedTermSet	Glossary pages	Definitional content is disproportionately cited
BreadcrumbList	Nested pages	Structure comprehension
HowTo	Process pages	Step extraction
Review / AggregateRating	Where genuine	Trust signal — only if real

Validate everything with Google's Rich Results Test. Never mark up content that is not visible on the page.

Content structure for extraction

The single highest-leverage on-site change:

Answer-first. After every H2, the very next thing is a direct, self-contained, 40–60 word answer to that heading. Detail follows after.

Also required:

- Headings shaped as real buyer questions, taken from the prompt set
- Comparison tables as real tables
- Every claim self-contained — an engine may lift one paragraph with no context
- Statistics with visible source attribution
- `dateModified` surfaced visibly; stale content drops out of synthesis
- Glossary and definitional pages — LLMs cite clean definitions heavily

Off-site levers, ranked by leverage

1. **Third-party "best X for Y" listicles.** The highest-leverage lever, because this is precisely what engines pull from for recommendation queries.
 2. **Original data or statistics.** One small original dataset earns more citations than fifty blog posts. Nothing is more citable than a number nobody else has.
 3. **Review platforms and credible industry directories.**
 4. **Editorial mentions in trusted publications.**
 5. **Entity consistency** — identical name, category, and description everywhere.
 6. **Wikipedia / Wikidata** — measurably significant (28 vs 52 days to first ChatGPT citation) **but only where the business genuinely meets notability requirements.** If it does not, the entry will be removed and the relationship damaged. Never sell this as a shortcut.
-

8. The prompt set — the core instrument

Everything measurable in this business comes from here. Build it first, build it properly, and then never change it mid-engagement.

Why it must be fixed

Baseline comparison is the entire product. If the prompt set changes, month three cannot be compared to month zero, and there is no proof of anything. Version it. Additions go into a new version, documented, and reported separately.

Composition — 40 to 50 prompts

Five categories. Proportions are a starting point, adjust per vertical.

Category	Share	Example shape
Category discovery	~30%	"best [service] in [city]"
Comparison	~20%	"[option A] vs [option B] cost and results"
Problem-led	~25%	"how do I fix [problem the service solves]"
Brand-direct	~15%	"is [brand] any good", "[brand] reviews"
Local / qualified	~10%	"[service] near [landmark]", "affordable [service] [city]"

Build these from how buyers actually speak, not from keyword tools. Keyword tools give you search syntax; AI prompts are conversational sentences.

Run methodology — conditions must be identical every time

Non-negotiable, or the data is worthless:

1. **Three runs per prompt per engine.** LLM outputs are non-deterministic. One run tells you nothing.
2. **Clean session.** Logged out, no history, no personalisation. A fresh browser profile or incognito.
3. **Consistent geography.** Same location every run. Note it. Use a consistent VPN exit if targeting a market you are not in.
4. **Same day of week, similar time of day.**
5. **Record verbatim.** Capture the actual answer text and the cited URLs, not a summary.
6. **Screenshot every result where the client is absent and a competitor is present.** These are your sales assets and your report evidence.

What to record per run

Field	Values
Prompt ID	P01–P50
Prompt text	verbatim
Engine	ChatGPT / Perplexity / Gemini / AI Overviews / Claude
Run number	1, 2, 3
Date, time, location	—
Client mentioned	yes / no
Client cited with link	yes / no
Position in answer	1st / 2nd / 3rd / later / absent
Competitors named	list, in order
Cited domains	full list of URLs
Screenshot reference	filename

The two headline metrics

Mention rate = (runs where client is named) ÷ (total runs) × 100 Per engine. Reported as a percentage.

Citation share = (runs where client is cited with a link) ÷ (total runs) × 100 Per engine. Stricter and more valuable than mention rate.

Also report **competitive position**: for each competitor, their mention rate alongside the client's. This is what makes the report land emotionally.

Worked example

40 prompts × 3 runs × 5 engines = **600 data points** per full cycle.

If the client is named in 24 of 120 ChatGPT runs: mention rate on ChatGPT = $24 \div 120 = 20\%$.

If cited with a link in 9 of those: citation share on ChatGPT = $9 \div 120 = 7.5\%$.

Report each engine separately. Never average across engines — the ~11% overlap makes a blended number meaningless.

Time cost — know this before pricing

Task	Estimated time
Build prompt set for a new vertical	2–3 h (reusable afterwards)
Adapt existing set to a new client	30–45 min
Run 600 data points	4–6 h
Record and clean data	1–2 h
Analysis and diagnosis	3–4 h
Write the report	3–4 h
Total first audit	13–19 h
Monthly re-run + report	6–9 h

These are estimates from task decomposition, not measured averages. Track your actual hours on the first three clients and replace them.

At \$1,500 for 13–19 hours, the effective rate is roughly \$80–115/hour. That is sound. Do not discount below the founding-client rate.

Tooling

Start entirely manual. Browser plus a spreadsheet. Cost: zero.

Do this for the first three clients even though it is tedious. Reasons: you learn what the data actually looks like, you can explain every number to a client, and you will be able to tell when a paid tool is wrong.

Consider a dedicated AI-visibility tracking tool only at 4+ clients, when the manual run becomes the bottleneck. Buying a tool before having clients is the most common way to lose money in this business.

9. Delivery — phase by phase, step by step

Five phases. Each broken into tasks with time estimates and a definition of done.

Phase 0 — Baseline (Days 1–3, the paid audit)

Purpose: establish the measurement that everything later is compared against. Skipping or rushing this destroys your ability to prove results at month three.

Step	Task	Time	Done when
0.1	Kickoff call: confirm services, competitors (max 3), target geography, target customer	45 min	Written and confirmed by client
0.2	Build or adapt the prompt set	30 min – 3 h	40–50 prompts, categorised, versioned
0.3	Client signs off on the prompt set	—	Written approval received
0.4	Run the full set: 3 runs × all engines	4–6 h	All data points recorded
0.5	Screenshot every gap (client absent, competitor present)	during 0.4	Filed and referenced
0.6	Technical audit: robots.txt, rendering, schema, sitemap, canonicals, noindex	1–2 h	Findings documented
0.7	Extractability audit: heading structure, answer-first, tables, self-containment	1 h	Page-level findings
0.8	Off-site audit: entity consistency, directories, reviews, third-party mentions, Wikipedia status	1–2 h	Gap list
0.9	Competitor analysis: which domains are being cited, and why	1–2 h	Source list with reasons
0.10	Score and prioritise fixes by expected movement vs effort	1 h	Ranked roadmap
0.11	Write the report	3–4 h	Report complete per Section 10
0.12	Delivery call: walk through findings and roadmap	45 min	Client understands next step

Step 0.3 matters. Prompt-set sign-off is what makes the baseline defensible. Without it, a client can later claim you measured the wrong things.

Phase 1 — Retrieval unblock (Week 1 of implementation)

Fastest ROI in the entire engagement. These changes propagate in days.

Step	Task	Time	Done when
1.1	Update <code>robots.txt</code> for all AI crawlers (Section 7 list)	30 min	Live and verified
1.2	Author and deploy <code>llms.txt</code>	1 h	Live at root
1.3	Fix rendering blockers: JS-only content, cookie walls, interstitials	1–4 h	Content present in raw HTML
1.4	Fix accidental <code>noindex</code> , broken canonicals, sitemap errors	1 h	Verified clean
1.5	Implement core schema: Organization/LocalBusiness, Service, FAQPage, BreadcrumbList	2–3 h	Passes Rich Results Test
1.6	Entity consistency pass: name, category, description identical everywhere	1–2 h	Audit sheet complete
1.7	Set up server log monitoring for AI crawler hits	1 h	Baseline captured, daily tracking live

Step 1.7 is your week-one proof. Capture logs *before* 1.1 so you have a before-and-after. "These crawlers weren't reaching you; now GPTBot hits daily" is verifiable and lands well.

Note on access: if the client's site is on WordPress, Wix, Webflow or similar, you can usually do this yourself with admin access. On a custom build without code access, you write precise instructions for their developer. You are the architect; you do not have to be the only pair of hands.

Phase 2 — On-site AEO (Weeks 2–4)

Step	Task	Time	Done when
2.1	Map prompts to target pages; identify gaps needing new pages	2 h	Mapping document
2.2	Rewrite money pages answer-first	6–10 h	Every H2 followed by a 40–60 word direct answer
2.3	Reshape headings into buyer-question form using prompt-set language	2 h	Deployed
2.4	Build FAQ blocks + FAQPage schema from real prompts	3 h	Live and validated
2.5	Convert comparison content to real <code><table></code> with <code><th scope></code>	1–2 h	Semantic markup verified
2.6	Create glossary / definitional pages with <code>DefinedTerm</code>	3–4 h	Minimum 10 terms live
2.7	Add visible source attribution to every statistic	1 h	All claims sourced
2.8	Surface <code>dateModified</code> on content pages	30 min	Visible

Targeting strategy: attack **low-competition long-tail prompts first**, not head terms. Those are what move in 4–12 weeks and become your month-two proof. Head terms are a month-six conversation.

AI usage rule for step 2.2: draft with AI, then verify every factual and service claim yourself. In healthcare, legal, and financial verticals an AI-invented claim is a liability for the client and for you. You are the reviewer. Never publish AI output unchecked.

Phase 3 — Off-site GEO (Week 3 onward, runs for months)

Start this in parallel with Phase 2. Do not wait for on-site completion.

Step	Task	Time	Cadence
3.1	Claim and correct all directory and platform listings	2–3 h	One-off, then monitor
3.2	Review platform strategy: which platforms, how to earn reviews ethically	2 h setup	Ongoing
3.3	Identify "best X for Y" listicles ranking for target prompts	2 h	Monthly refresh
3.4	Outreach for inclusion in those listicles	3–5 h	Ongoing
3.5	Produce an original data asset with the client	5–10 h	Once per quarter
3.6	Digital PR / editorial placement outreach	ongoing	Monthly
3.7	Wikipedia/Wikidata assessment — only if notability is genuinely met	2 h assessment	Once

Step 3.5 is the highest-leverage item in this phase. A small original dataset — "average treatment timeline across 500 patients", "our fee data across 200 cases" — is more citable than any amount of written content. Engines cite numbers nobody else has. Push clients toward this.

Step 3.7 warning: Wikipedia has strict notability requirements and active editors. Creating an entry for a business that does not meet them will get it removed and can damage the client's reputation. Assess honestly; decline if it does not qualify.

Phase 4 — Monitoring (Month 2 onward)

Step	Task	Time	Cadence
4.1	Re-run the full prompt set, identical conditions	4–6 h	Monthly
4.2	Update scorecard, compute deltas vs baseline and vs last cycle	1 h	Monthly
4.3	Competitor movement analysis	1 h	Monthly
4.4	Crawler log review	30 min	Monthly
4.5	Produce the report	2–3 h	Monthly
4.6	Report call	45 min	Monthly
4.7	Content freshness pass — visibility decays	2 h	Quarterly
4.8	Prompt set review — add new prompts as a new version	1 h	Quarterly

Full engagement summary

Phase	Duration	Your hours
0 — Baseline audit	2 weeks	13–19 h
1 — Retrieval unblock	1 week	7–12 h
2 — On-site AEO	3 weeks	19–28 h
3 — Off-site GEO	ongoing	15–25 h initial, then 8–12 h/mo
4 — Monitoring	ongoing	9–13 h/mo

Month 1 of a full engagement: roughly 40–60 hours. Steady state from month 3: roughly 12–20 hours per client per month.

Estimates from task decomposition. Measure your real hours on client one and replace these numbers.

10. Reporting

The report is the product the client actually holds. It justifies the retainer.

Audit report structure

1. **Executive summary** — one page. Where they stand, the three biggest gaps, what to do first. Written for someone who reads nothing else.
2. **Baseline scorecard** — mention rate and citation share, per engine, in a table. Numbers, not adjectives.
3. **Competitive position** — the client versus up to three named competitors, per engine. The most persuasive page in the document.
4. **Evidence** — the screenshots. Client absent, competitor present, prompt shown.
5. **Diagnosis** — organised by the three requirements from Section 7: access, extractability, trust. Each finding states what is wrong, why it matters, and what fixes it.
6. **Prioritised roadmap** — ranked by expected movement against effort. Marked as on-site or off-site, with an expected timeframe drawn from Section 6.
7. **Methodology appendix** — the prompt set, run conditions, dates, definitions of every metric. This is what separates a real audit from a generated one.

Monthly report structure

Shorter. Four pages maximum.

1. **Movement summary** — mention rate and citation share this cycle versus baseline versus last cycle, per engine.
2. **What moved and why** — tie changes to specific work completed.
3. **Leading indicators** — crawler activity, placements secured, schema status, snippets gained. Especially important in months 1–3 when citations are still thin.
4. **Next cycle's priorities** — the fix list, ranked.

Reporting rules

- **Report bad months honestly.** Non-deterministic systems fluctuate. A dip explained holds a client; a dip hidden loses them permanently when discovered.
 - **Never average across engines.** Report each separately.
 - **Always show the comparison** — versus baseline and versus named competitors. Absolute numbers alone mean nothing to a client.
 - **Attribute honestly.** If a change cannot be tied to your work, say so.
 - Same format every month. Consistency is itself a credibility signal.
-

11. Client acquisition

Channel priority

Ranked for a solo operator with no case studies. Spend roughly 70% of sales time on channel 1.

#	Channel	Why	Speed to first client
1	Direct outbound with a personalised audit finding	You arrive with evidence, not a pitch	2–4 weeks
2	Agency / web-developer partnerships	They hold clients and lack this capability. Recurring, zero marketing.	3–6 weeks, then compounding
3	Niche communities where owners already gather	Value first, then inbound DMs	3–6 weeks
4	LinkedIn content — anonymised audit findings	Builds inbound over time	6–12 weeks
5	Freelance platforms — specific keywords only	Buyers ready to pay, but you look like a commodity	Variable. Backup only.

Not now: paid ads (no conversion data yet), cold calling without proof, generic "AI agency" broadcasting, presence on every platform.

Channel 1 — the core motion

The entire advantage is that you show up with proof of *their* specific gap.

Step 1 — Build the list (30–45 min for 50 prospects). Google Maps, industry directories, or LinkedIn for the chosen vertical and city. Capture: business name, website, owner name, email, LinkedIn URL.

Step 2 — Qualify (5 min each, Section 4 Stage 1). Score all 50. Keep those scoring 4+. Expect roughly 20–30 to survive.

Step 3 — Run the mini-audit (3–4 h for the surviving list). Five core prompts per prospect across ChatGPT, Perplexity and AI Overviews. Record: is the prospect named? which competitors are? Screenshot every gap.

Step 4 — Personalised outreach. One email or DM per prospect, each referencing their own gap and attaching their own screenshot.

Step 5 — Follow up. Two follow-ups. Most replies come from these, not the first message.

Outreach templates

Adapt every one. Never send generic.

First touch — email

Subject: [Competitor] comes up in ChatGPT, [Their business] doesn't

Hi [Name],

I ran a few searches the way your patients now do — asking ChatGPT and Perplexity instead of Google.

For "[actual prompt]", the answer named [Competitor A] and [Competitor B]. [Their business] wasn't mentioned. Screenshot attached.

These are people actively choosing a [service] provider, and they're choosing without seeing you. It's fixable — the reason is usually technical, not reputational.

Worth 15 minutes to walk you through what I found and what causes it?

[Name] · benamedfirst.com

Why this works: specific, evidenced, no adjectives, no pitch, one small ask.

Follow-up 1 — 4 days later

Hi [Name], following up on the ChatGPT search I sent.

One thing worth adding: I checked which sources those engines *are* pulling from for [service] in [city]. Three of them are directories you're either not listed on or listed inconsistently. That's a large part of the gap and it's a straightforward fix.

Happy to send the list whether or not we ever work together — want it?

Gives value unconditionally. Frequently the reply-generating message.

Follow-up 2 — 8 days after that, closing the loop

Hi [Name], I'll stop here so I'm not cluttering your inbox.

If it's useful later, the summary is: for [prompt], AI engines currently name [Competitor A] and [Competitor B] and not you. My notes are attached — yours to use with anyone.

If it becomes a priority, you know where I am.

Non-needy, leaves the door open, leaves value behind. A meaningful share of replies come from this one.

LinkedIn DM — shorter

Hi [Name] — I ran "[prompt]" through ChatGPT and Perplexity. Both named [Competitor] and not [Business]. Happy to send the screenshots and the reason — want them?

Agency partnership email

Subject: The AI-visibility question your clients are starting to ask

Hi [Name],

Your clients have probably started asking why they don't appear when someone asks ChatGPT for a recommendation. It's a different discipline from SEO — different signals, different measurement, no Search Console equivalent.

I do this work white-label for agencies. You keep the relationship and the margin; I deliver measurement and implementation under your brand.

I've run a sample check on one of your clients — [Client] doesn't appear for [prompt] while two competitors do. Screenshot attached, no charge.

Worth a conversation?

Volume mathematics

Estimates from typical proof-led outbound response patterns. Track your own and replace them.

Stage	Rate	From 50 qualified prospects
Qualified after Stage 1	~50% of raw list	25
Reply to personalised, evidenced outreach	15–25%	4–6
Reply → call booked	~60%	3–4
Call → audit sold	30–50%	1–2

Roughly 40–50 personalised touches produces 1–2 audits.

Complete the full set before judging the channel. Sending five emails and concluding "it doesn't work" is the most common failure mode.

The specificity multiplier

The largest lever is not the channel — it is specificity.

Message	Reply rate
"I do AI SEO for businesses"	~1%
"Your clinic is invisible in ChatGPT — here's the screenshot"	15–20%

Same channel, roughly 15x difference. Every channel follows this rule: never go generic, always lead with their own gap.

Email deliverability — do not skip

Get this wrong and outreach fails silently. No bounce, no error, just no replies.

- Never send cold outreach from personal Gmail or Outlook
- Use a domain mailbox: `you@benamedfirst.com` (Google Workspace or Zoho)

- Configure **SPF, DKIM and DMARC** DNS records before the first send
- Warm the domain: 10–20 sends/day in week one, increase gradually. A new domain blasting 100 emails on day one gets burned
- Test with a deliverability checker before any real send
- Plain text, one link maximum, no tracking pixels, no images in the first email

Tracking

From day one, record: touches sent, replies, calls booked, proposals sent, deals closed. Per channel, per vertical.

Once you know that 45 touches produces one client, "no work coming in" stops being panic and becomes arithmetic: send 45 more.

12. Sales and closing

Call structure — 30 minutes

Minutes	Purpose	What you do
0–3	Frame	"I'll show you what I found, then you decide if it's worth pursuing."
3–10	Show the evidence	Share screen. Run the prompt live if you can. Let them see competitors named and themselves absent. Then stop talking.
10–18	Qualify	The six questions from Section 4 Stage 2
18–24	Explain and scope	What causes it, what the audit produces, honest timeline from Section 6
24–28	Price and terms	The number, what is included, what is not
28–30	Close	One clear next step

The silence at minute 10 is the most important part of the call. Show the evidence and stop. Let them react. Their reaction tells you whether this closes.

Running the demo live

If possible, run a prompt live on the call rather than showing a screenshot. It is dramatically more persuasive because it cannot be doubted.

Prepare: clean browser profile, logged out, correct location, and two or three prompts you have already verified show the gap. Have screenshots as backup in case of a bad connection or an unexpectedly different answer.

Objection handling

"How do I know this works?"

You don't yet, and I won't promise it. What I can promise is measurement. We establish exactly where you stand today, then report the same numbers monthly. If nothing moves, you'll see that clearly. Most vendors in this space can't tell you either way — that's the actual problem.

"Can you guarantee I'll show up in ChatGPT?"

No, and anyone who does is lying to you. These are non-deterministic systems — the same question can return different answers on consecutive runs. What I control is whether the engines can access your content, whether it's structured to be extracted, and whether credible sources reference you. Those are the inputs. I report the outputs honestly.

"Why not just do SEO?"

You should still do SEO. But Google rankings and AI citations don't map one-to-one — I can show you queries where you rank on page one and still don't appear in the AI answer. Different mechanism, different work.

"This is expensive."

What's a new [patient/client] worth to you? [Listen.] So the audit costs less than [fraction] of one. And you keep the roadmap whether or not you continue with me.

"How long until results?"

Crawler access changes within days and I can show you that in your server logs in week one. First citations typically 4 to 12 weeks on lower-competition queries. Consistent visibility across engines, 3 to 6 months. Anyone promising faster is selling you something.

"Let me think about it."

Of course. What specifically would you want to be more confident about? [Then address only that.]

"Send me a proposal."

Happy to. Before I do — is there anything in what I've shown you that doesn't land yet? I'd rather resolve it now than in a document.

"Can we do results-based pricing?"

No. With non-deterministic systems, outcomes aren't fully within anyone's control, and results-based pricing would push me toward whatever moves a number fastest rather than what's actually right for your business. I charge for the work and I report the results honestly.

The proposal

One page. Send within 24 hours. Long proposals lose deals.

Contents: what you found (2–3 lines), scope, deliverables, timeline, price, payment terms, what is explicitly excluded, start date, one clear acceptance instruction.

Terms to hold

- **Audit: 100% upfront.** Standard for fixed-scope diagnostics.
- **Implementation: 50/50.**
- **Retainer: monthly in advance, 3-month minimum.**
- No results-based pricing. No guarantees. No free pilots.
- One round of report revisions included; further rounds are billable.
- Scope changes are quoted separately in writing.

Payment mechanics — set up before the first close

Nothing kills momentum like a client saying yes and then waiting three days for your bank details.

- Wise Business or Payoneer for international transfers; Stripe for card payments where available
 - An invoice template ready: your details, client details, scope line, amount, currency, due date, payment instructions
 - **Consult a CA about export-of-services rules and GST/LUT for foreign clients before your first invoice.** This is a genuine compliance area and getting it wrong is expensive to unwind. This handbook does not give tax advice.
-

13. Targets and unit economics

All figures below are **planning estimates**, built from the funnel rates in Section 11 and the hour estimates in Section 9. They are not forecasts and not promises. Replace them with measured numbers after client three.

Capacity ceiling

This is the real constraint, not lead flow.

	Month 1 of engagement	Steady state (month 3+)
Hours per client	40–60	12–20
Clients at 100 h/month	~2 new	~5–7 ongoing
Clients at 140 h/month	~3 new	~7–10 ongoing

Practical ceiling for one person: about 5–6 retained clients plus one new onboarding at a time. Beyond that, either delegate (Section 15) or raise prices.

Note the shape: onboarding is 3–4x more expensive in hours than steady state. Two new clients in the same month is a bad month. Stagger starts deliberately.

Ramp plan — first six months

Deliberately conservative. Hitting these means the business works.

Month	Focus	Audits	Implementations	Retainers	Est. revenue
1	Setup, prompt set, 50 touches	1–2 @ \$750	0	0	\$750–1,500
2	2 case studies, 50 more touches	1–2 @ \$750	1 @ \$5,000	0	\$5,750–6,500
3	Price to \$1,500, first retainers	2 @ \$1,500	1 @ \$5,000	1 @ \$1,500	\$9,500
4	Agency partnerships open	2 @ \$1,500	1 @ \$5,000	2 @ \$1,500	\$11,000
5	Referrals begin	2 @ \$1,500	2 @ \$5,000	3 @ \$1,500–3,000	\$17,500+
6	Steady state	2 @ \$1,500	1–2 @ \$5,000	4–5 @ \$1,500–3,000	\$19,000+

Deal targets: 1–2 audits per month in months 1–2, then 2 per month sustained.

The retainer line is what matters. Audits and implementations are one-time; retainers compound. By month 6, retainers should be your revenue floor.

Outreach volume required

Working backwards from ~45 personalised touches per client:

Target	Touches needed	Per week
1 audit/month	~45	~11
2 audits/month	~90	~23
3 audits/month	~135	~34

At roughly 8 minutes per personalised touch including the mini-audit, 90 touches is about **12 hours per month** of prospecting. Budget it as a fixed weekly block, not something done when work is thin.

The most common failure in this business is stopping outreach after landing two clients, then having an empty pipeline in month four. Prospecting is a permanent weekly commitment.

Cost base

Item	Monthly
Domain (benamedfirst.com, amortised)	~\$1
Email (Zoho or Google Workspace)	\$3–7
Claude / ChatGPT subscription	\$20
VPN (consistent geo for prompt runs)	\$5–12
Site hosting (Vercel free tier viable)	\$0–20
Total to start	\$30–60

Deferred until 4+ clients: AI-visibility tracking tool (\$50–200/mo), SEO tooling, VA support.

Operating costs are near zero. Your constraint is hours, not money. This is why capacity planning matters more than cost control.

Break-even

At \$30–60/month of cost, a single \$750 founding audit covers more than a year of operating expenses. There is no meaningful financial risk in starting. The risk is entirely in time spent on the wrong activities.

Revenue mix target by month 6

Source	Share	Why
Retainers	50–60%	Predictable, compounding
Implementation	30–40%	Higher ticket, lumpy
Audits	10–20%	Entry point, not the business

If audits are still the majority of revenue at month 6, the ladder is not converting. Investigate why audits are not becoming implementations — usually the roadmap is not concrete enough, or timeline expectations were mishandled.

Leading indicators to watch weekly

1. Touches sent
2. Reply rate
3. Calls booked
4. Audit-to-implementation conversion
5. Retainer count and churn
6. Hours per client (against Section 9 estimates)

Item 6 is the one people ignore. If a client consumes double the estimate, either the scope grew or the price is wrong. Margin erosion from scope creep is the documented reason most agencies in this space fail — not lack of demand. [Angel Investors Network]

14. Risk register

Risk	Likelihood	Impact	Mitigation
Client churns at month 2 before citations move	High	High	Leading indicators weekly from week 1 (Section 6). 3-month minimum term.
Scope creep destroys margin	High	High	Written scope. Changes quoted separately. Track hours per client.
Client demands a guarantee	Medium	Medium	Disqualify early (Section 4). Scripted response (Section 12).
Overpromised timeline leads to refund request	Medium	High	Present the Section 6 timeline table on every sales call, before pricing.
Cold email lands in spam	High if unaddressed	High	SPF/DKIM/DMARC + domain warming (Section 11).
Cannot modify client's site	Medium	High	Stage 1 check 4. Confirmed again on the call.
Engine changes behaviour and results reverse	Medium	Medium	Never promise outcomes. Report honestly. Diversify across engines.
Prompt set changed mid-engagement, baseline lost	Low	Severe	Version control. New prompts go in a new version, reported separately.
Two clients onboard in the same month, delivery slips	Medium	High	Stagger start dates deliberately.
Pipeline empties in month 4	High	High	Fixed weekly prospecting block, regardless of current workload.
AI-drafted content contains a false claim	Medium	Severe	You verify every factual claim. Non-negotiable in health, legal, finance.
Wikipedia entry removed, client embarrassed	Low	High	Assess notability honestly. Decline where it does not qualify.
Tax/compliance error on foreign invoices	Medium	Medium	Consult a CA before the first invoice.

15. Team, tools and delegation

Do it all yourself for the first three clients

Not for cost reasons. For three specific ones:

1. You learn the **actual** time each task takes, which is the only sound basis for pricing.
2. You learn what can safely be delegated and what cannot.
3. You stay cash-flow positive with no fixed commitments.

Hiring before you have measured your own delivery is the fastest way to lose money here.

Work composition

Roughly, across a full engagement:

Type	Share	Notes
Judgement and manual work	~60%	Measurement, verification, client calls, strategy. This is what you are paid for. AI cannot do it.
AI-accelerated	~30%	Content drafts, schema code, report summaries. You review everything.
Delegable grunt work	~10%	Prompt running, data entry, listing cleanup

When to delegate, and what

Only after 4–5 clients, and task by task — not by hiring a team.

Delegate first: prompt running and data entry. A VA can execute a documented run methodology and fill the scorecard. Biggest hour saving available.

Delegate later: content drafting (per-article freelancer), off-site outreach, directory cleanup.

Never delegate: measurement methodology, technical and schema work, client communication, final quality review. These are the moat.

Structure: a VA plus per-task freelancers. No fixed salaries. Variable cost only when work exists.

Tool stack

From day one:

Tool	Purpose	Cost
Browser + spreadsheet	Measurement	Free
Claude / ChatGPT paid	Content and code assistance	\$20/mo
Google Rich Results Test	Schema validation	Free
Screaming Frog (free tier)	Technical crawl, up to 500 URLs	Free
Deliverability checker	Email testing	Free
VPN	Consistent geography for runs	\$5–12/mo
Calendar booking tool	Call scheduling	Free tier

Only at 4+ clients: a dedicated AI-visibility tracking platform. Run manually first so you can tell when a tool is wrong, and so you can explain every number to a client.

16. First 30 days

Week 1 — Foundation

Day	Task	Output
1	Buy benamedfirst.com. Set up domain email. Configure SPF/DKIM/DMARC.	Sending infrastructure live
1	Choose one vertical and one geography	Written decision
2	Build the prompt set for that vertical (40–50 prompts, 5 categories)	Versioned prompt set
3	Build the scorecard template	Reusable spreadsheet
4	Run the full methodology on your own site	Workflow learned, before/after material
5	Set up LinkedIn profile: headline as offer, banner, about section	Profile live

Week 2 — Proof

Day	Task	Output
6–7	Build a 50-prospect list, qualify with Stage 1	~25 qualified prospects
8–9	Run mini-audits on all qualified prospects, screenshot every gap	Evidence library
10	Pick the 5 clearest gaps, produce full sample audit reports	Portfolio deliverables

Week 3 — Outreach

Day	Task	Output
11	Send first 10 personalised emails	Sending begins, domain warming
12–15	15–20 touches/day, ramping. Follow-ups on schedule.	60–80 touches sent
12–15	Post 2–3 LinkedIn posts using anonymised findings	Profile is not dead

Week 4 — Close

Day	Task	Output
16–20	Take calls. Follow the Section 12 structure.	2–4 calls
16–20	Continue outreach — do not stop because calls are happening	Pipeline maintained
21+	Send proposals within 24 hours. Close 1–2 founding audits.	First revenue

The 30-day success criterion

One paid audit, and a pipeline that is still full.

Not ten clients. Not a perfect website. One paying client and continuing momentum.

What to explicitly not do in the first 30 days

- Build a logo system or brand guidelines beyond a wordmark and favicon
- Register a company
- Buy any paid tool beyond an AI subscription
- Build a multi-page site before the prompt set exists
- Spend more than 30 minutes on any naming or branding decision
- Learn another skill instead of doing outreach

Every one of these feels productive and none of them produces a client.

17. Integrity rules — binding

These are not guidelines. Breaking any one of them creates legal exposure, destroys client trust, or both. They apply to everything: the website, proposals, reports, outreach, social content, and anything an AI assistant produces from this handbook.

Never fabricate

- No invented client names, logos, testimonials, or case study numbers
- No invented statistics. If a figure is not in Section 18, do not state it as fact.
- No claimed awards, certifications, partnerships, or years of experience
- No implying a team exists. This is a solo practice — position it as a specialist practice, not a fake agency.
- No client names in examples without written permission. Anonymise as "a [vertical] business in [region]".
- No trademark symbols. The name is not registered.

On the website, build components for testimonials and case studies, populate them with data explicitly flagged as placeholder, and do not publish until real ones exist. A client can verify fabricated proof in minutes.

Never guarantee

No guarantees of rankings, positions, citations, mention rates, timelines for specific outcomes, or revenue impact. LLM outputs are non-deterministic; any guarantee would be knowingly false.

Sell the process and honest measurement.

Always attribute

Every statistic used publicly carries its source. Section 18 is the approved list. When paraphrasing a source, do not distort its meaning.

Verify AI output

AI drafts content; you verify every factual and service claim before it is published or sent. In healthcare, legal, and financial verticals a hallucinated claim can create genuine liability for the client and for you. You are the reviewer of record.

Respect the platforms

- Follow `robots.txt` conventions and platform terms
- No fake reviews, no incentivised reviews where prohibited, no review gating
- No Wikipedia entries for businesses that do not meet notability requirements
- No manipulation of engagement metrics

Distinguish estimate from fact

This handbook contains both. Estimates — hour figures, funnel rates, revenue ramp — are labelled. Never present an estimate to a client as a measured result.

Handle client data carefully

Server logs, analytics access, and CMS credentials are sensitive. Store credentials in a password manager, never in plain text or in chat. Never echo credentials into a document or an AI prompt.

18. Sources

Every statistic in this handbook, with its origin. Verify before republishing — this space changes fast.

Market size and demand

- **[3way]** AI search visits +42.8% YoY vs Google +2.4%; AI-referred visitor value ~4.4x — 3way.social/blog/how-to-make-money-with-aeo-geo
- **[Conductor]** 97% reported positive AEO impact; 94% increasing 2026 investment — conductor.com/academy/state-of-aeo-geo-report
- **[LLMPulse]** "geo agency" searches +2,300% YoY — llmpulse.ai/blog/geo-agency-guide
- **[Algomizer]** GEO market USD 1.09B (2026) → 17.1B (2034), 40.6% CAGR — algomizer.com/blog/generative-engine-optimization-services
- **[SomethingInc]** Cited links in AI overviews earn 8–12% CTR (Princeton / Georgia Tech research) — somethinginc.com/blog/how-generative-search-engines-cite-brands

Pricing benchmarks

- **[OutreachBloom]** AEO/GEO \$2,000–10,000/mo; audits \$1,000–4,000; enterprise \$25,000+ — outreachbloom.com/aeo-geo-pricing-guide
- **[Digital Elevator]** Mid-market retainers \$2,000–10,000/mo — thedigitalelevator.com/blog/aeo-and-geo-pricing-guide
- **[Fuel Online]** Full AEO/GEO programmes \$8,000–15,000/mo — fuelonline.com/ai-seo-geo/ai-seo-pricing-aeo-geo-cost-2026

Timelines

- **[LeanLabs]** Crawler access fixable in days; first citations 4–12 weeks on lower-competition prompts — leanlabs.com/faq/answer-engine-optimization/how-long-does-aeo-take-to-show-results
- **[AirOps]** First mentions typically 30–60 days — airops.com/blog/results-from-aeo
- **[WebFX]** Early signals 4–8 weeks; stronger visibility 3–6 months — webfx.com/blog/ai/how-long-geo-results
- **[Upgrowth]** Initial citations 2–3 months; consistent visibility 3–6 months — upgrowth.in/how-long-does-geo-take
- **[Higoodie]** Weeks 1–2 first mentions, volatile, not yet wins — higoodie.com/blog/how-long-does-aeo-take
- **[CNABKE]** GEO requires 3–6 month validation cycle — cnabke.com/en/blogs/geo-validation-cycle-3-6-months.html

How the engines work

- **[FiveBlocks]** Retrieval engines reflect new sources in hours-to-weeks; training baselines update in months — fiveblocks.com/knowledge/ai-llm-reputation-management
- **[Broworks]** Perplexity real-time retrieval; ChatGPT blended; Claude training-anchored with Brave fallback — broworks.net/blog/how-do-perplexity-chatgpt-and-claude-choose-what-content-to-surface

- **[Topify]** ~11% cited-domain overlap across four engines; Wikipedia 7.8% of ChatGPT citations, ~48% of top-10 sources; first citation 28 days with Wikipedia vs 52 without; 80.41% of ChatGPT citations are .com — topify.ai/blog/llm-citation-patterns-ai-platforms
- **[Digital Applied]** 30-day log study: GPTBot ~4,200 hits/site/day, ClaudeBot ~1,800, PerplexityBot ~980; all respect robots.txt — digitalapplied.com/blog/agent-crawler-behavior-30-day-site-log-study
- **[Forrester]** Retrieval vs pre-training memory systems — duaneforresterdecodes.substack.com
- **[DataStudios]** Perplexity retrieval-augmented architecture — datastudios.org
- **[EverythingPR]** Five dominant answer engines; browsing on/off changes update speed — everything-pr.com

Definitions

- **[Adobe]** AEO vs GEO definitions — business.adobe.com/blog/aeo-vs-geo
- **[Stackmatix]** AEO = extracted answer; GEO = cited in synthesis; AIO is a Google feature, not a strategy — stackmatix.com/blog/aeo-vs-geo-difference
- **[SearchAtlas]** AEO governs delivery, GEO governs source selection — searchatlas.com/blog/aeo-aio-geo
- **[Contently]** All three converge on the same operational work — contently.com/2026/04/29/aeo-vs-geo-vs-llmo
- **[mo.agency]** Industry terminology is not settled — mo.agency/blog/seo-aeo-geo-explained

Why not web development

- **[TechRadar / 20i]** 75–76% of web design professionals report AI competition already affecting business — techradar.com/pro/website-hosting; 20i.com/blog/2026-web-designer-survey
- **[AInvest]** Namecheap 2026 survey, same signal — ainvest.com/news/10-000-website-dead-2606
- **[pikapod]** AI builders collapsed the professional site price floor — cryptic-caracal.pikapod.net/ai-website-killed-5000-homepage

Market structure

- **[Ciela — saturation]** Market not saturated, but the generic version is crowded — ciela.ai/blogs/is-ai-automation-agency-saturated-2026
- **[Ciela — failure]** Agencies fail competing on price, the one axis with no floor — ciela.ai/blogs/why-ai-automation-agencies-fail-data
- **[Ciela — verticals]** Same scope pays \$3,500 in one industry and \$35,000 in another — ciela.ai/blogs/ai-automation-niches-highest-paying
- **[Angel Investors Network]** Margin erosion from scope creep, not lack of demand, kills agencies — angelinvestorsnetwork.com/startups
- **[Goodspeed]** 79% started with AI agents, ~11% run them in production — goodspeed.studio/blog/ai-automation-agency
- **[Trillet / RAND]** 80%+ of AI projects fail, roughly double the non-AI IT rate — trillet.ai/blogs/is-the-ai-agency-bubble-real-what-the-data-says
- **[BCG]** \$200B tech-services opportunity; buyers want design + deploy + operate — bcg.com/publications/2026/the-200-billion-dollar-ai-opportunity-in-tech-services

LinkedIn specifications

- **[Contentin]** Personal banner 1584×396 — contentin.io/blog/linkedin-banner-size
- **[Neal Schaffer]** Company cover 4200×700, 6:1, max 3MB — nealschaffer.com/linkedin-banner-size
- **[RightBlogger]** The widely-cited 1128×191 company figure is outdated — rightblogger.com/blog/linkedin-banner-size
- **[Hyperclapper]** Personal banner 4:1, max 8MB — hyperclapper.com/blog-posts/linkedin-cover-photo-size-guide
- **[DoubleShot]** Centre 1200×300 safe zone to avoid mobile cropping — doubleshotcreative.com/blog/linkedin-image-sizes-cheat-sheet
- **[imResizer / okzest]** Profile photo 400×400; posts 1080×1080 or 1080×1350; link preview 1200×627 — imresizer.com/blog/linkedin-image-sizes-2026-complete-guide

Source content has been paraphrased and summarised throughout this handbook for compliance with licensing restrictions. Figures are reproduced as facts; wording is not.

19. Glossary

AEO — Answer Engine Optimization. Structuring content so an engine extracts it as a direct answer. Surfaces: featured snippets, People Also Ask, voice, AI Overviews. Mostly on-site.

AIO — AI Overviews. Google's AI-generated answer summaries. A platform feature, not a strategy. Anyone selling "AIO services" is describing a feature as a discipline.

Answer engine. A system that responds to a question with a synthesised answer rather than a list of links.

Answer-first writing. Placing a direct, self-contained 40–60 word answer immediately after a heading, before any elaboration. The highest-leverage on-site change for AEO.

Baseline. The measured starting position before any work. Without it, no result can be proven.

Citation share. Percentage of prompt runs in which the client is cited with a link. Stricter and more valuable than mention rate.

Entity consistency. Identical business name, category, and description across every property on the web. Helps engines resolve you as one entity.

Extraction. An engine lifting text from a page more or less verbatim. The AEO mechanism.

GEO — Generative Engine Optimization. Building the authority and presence that cause an LLM to cite you when synthesising an answer. Mostly off-site.

LLMO. Another term used roughly interchangeably with GEO. The terminology in this space is not settled.

llms.txt. A root-level plain-text file summarising a site for LLMs. Low cost, unproven impact. Do not oversell.

Maskable icon. An app icon with content inside a safe zone so Android can crop it to any shape without cutting the mark.

Mention rate. Percentage of prompt runs in which the client is named, cited or not.

Non-determinism. The property that the same prompt can produce different answers on consecutive runs. The reason guarantees are impossible and three runs per prompt are mandatory.

Prompt set. The fixed, versioned collection of 40–50 prompts used to measure visibility. The core instrument of the business.

RDAP. Registration Data Access Protocol — the official successor to WHOIS. Used by the domain-finder tool, not part of client work.

Retrieval layer. The live web search an engine performs at query time. Reachable, fast-moving, and the target of all our work.

Structured data / JSON-LD / schema. Machine-readable markup describing page content. Helps both extraction and entity resolution.

Synthesis. An engine blending multiple sources into one answer. The GEO mechanism.

Training layer. A model's baseline knowledge, fixed until retraining. Not reachable. Never promise anything about it.

20. Quick reference card

The pitch. We get businesses named and cited inside AI-generated answers, and we measure it.

The wedge. A screenshot of an AI engine naming their competitors and not them.

The differentiator. Measurement. There is no Search Console for AI answers.

Ideal client. New customer worth \$2,000+, already spending on marketing, competitor cited and they are not.

Best first market. UAE/GCC — high ticket, Western rates, ~1.5 hour timezone gap.

Prices. Audit \$1,500 fixed (\$750 founding). Implementation from \$5,000. Monitoring from \$1,500/mo, 3-month minimum.

The timeline sentence. "Month one is plumbing. Months two and three are when citations start. Months three to six are when visibility becomes consistent."

Never. Guarantee anything. Fabricate anything. Average across engines. Change the prompt set mid-engagement.

Always. Baseline before work. Three runs per prompt. Report per engine. Report bad months honestly.

Volume. ~45 personalised touches → 1 audit. 90 touches/month → 2 audits/month.

Capacity. 5–6 retained clients plus one onboarding, solo.

The weekly non-negotiable. A fixed prospecting block, regardless of how busy delivery is. Empty pipelines in month four come from skipping this in month two.

End of handbook. Version 1.0, August 2026.

Review and update quarterly. This market moves fast — timelines, engine behaviour, and pricing benchmarks all shift. Re-verify Section 18 before relying on any figure in a client-facing document.